Introduction to NLP

CS 4740 / CS 5740 / LING 4474 / COGST 4740

- Instructor: Claire Cardie
  - Professor in CS and IS (and CogSci)
- Three TAs
  - Ozan Irsoy
  - Tianze Shi
  - Sai Ram Sanapureddy
- One dog
  - Marseille (mahr-say)

Topics for Today

- What is NLP?
- What's involved? Why is it hard?
- Course structure and requirements

Introduction to NLP

CS 4740 / CS 5740 / LING 4474 / COGST 4740

- From Courses of Study

Computationally oriented introduction to natural language processing, the goal of which is to enable computers to use human languages as input, output, or both. Possible topics include parsing, ...

Natural Language Processing (NLP)

- "Natural" language
  - Languages that people use to communicate with one another
- Ultimate goal
  - To create computational models that perform as well at using natural language as humans do
- Immediate goal
  - To develop (and implement) algorithms that can process text ~~and speech~~ more intelligently

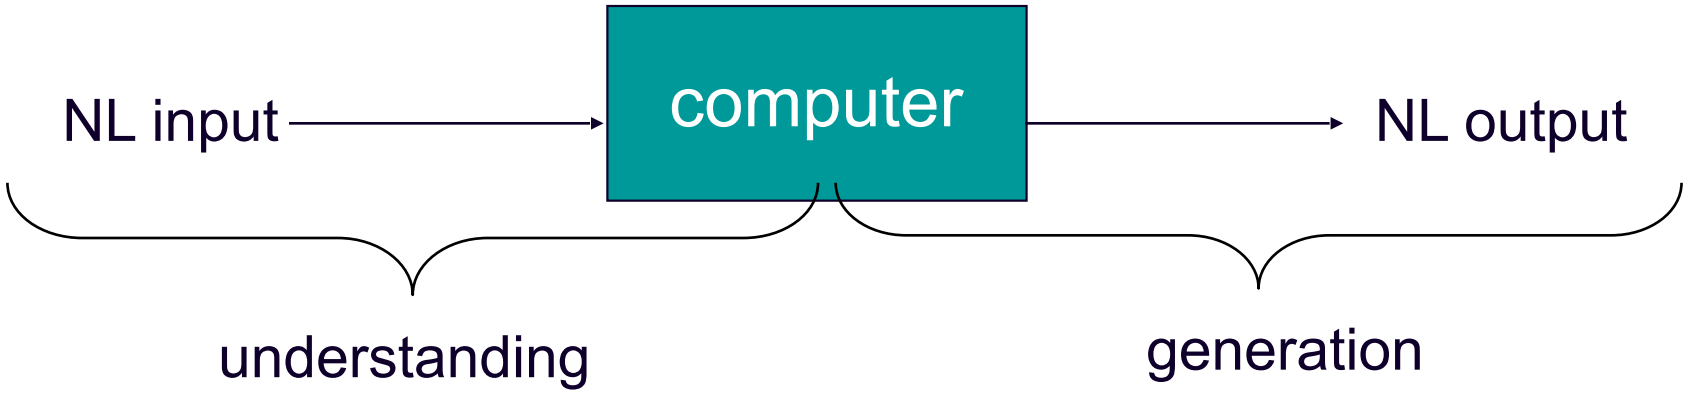

Information retrieval

- Web search

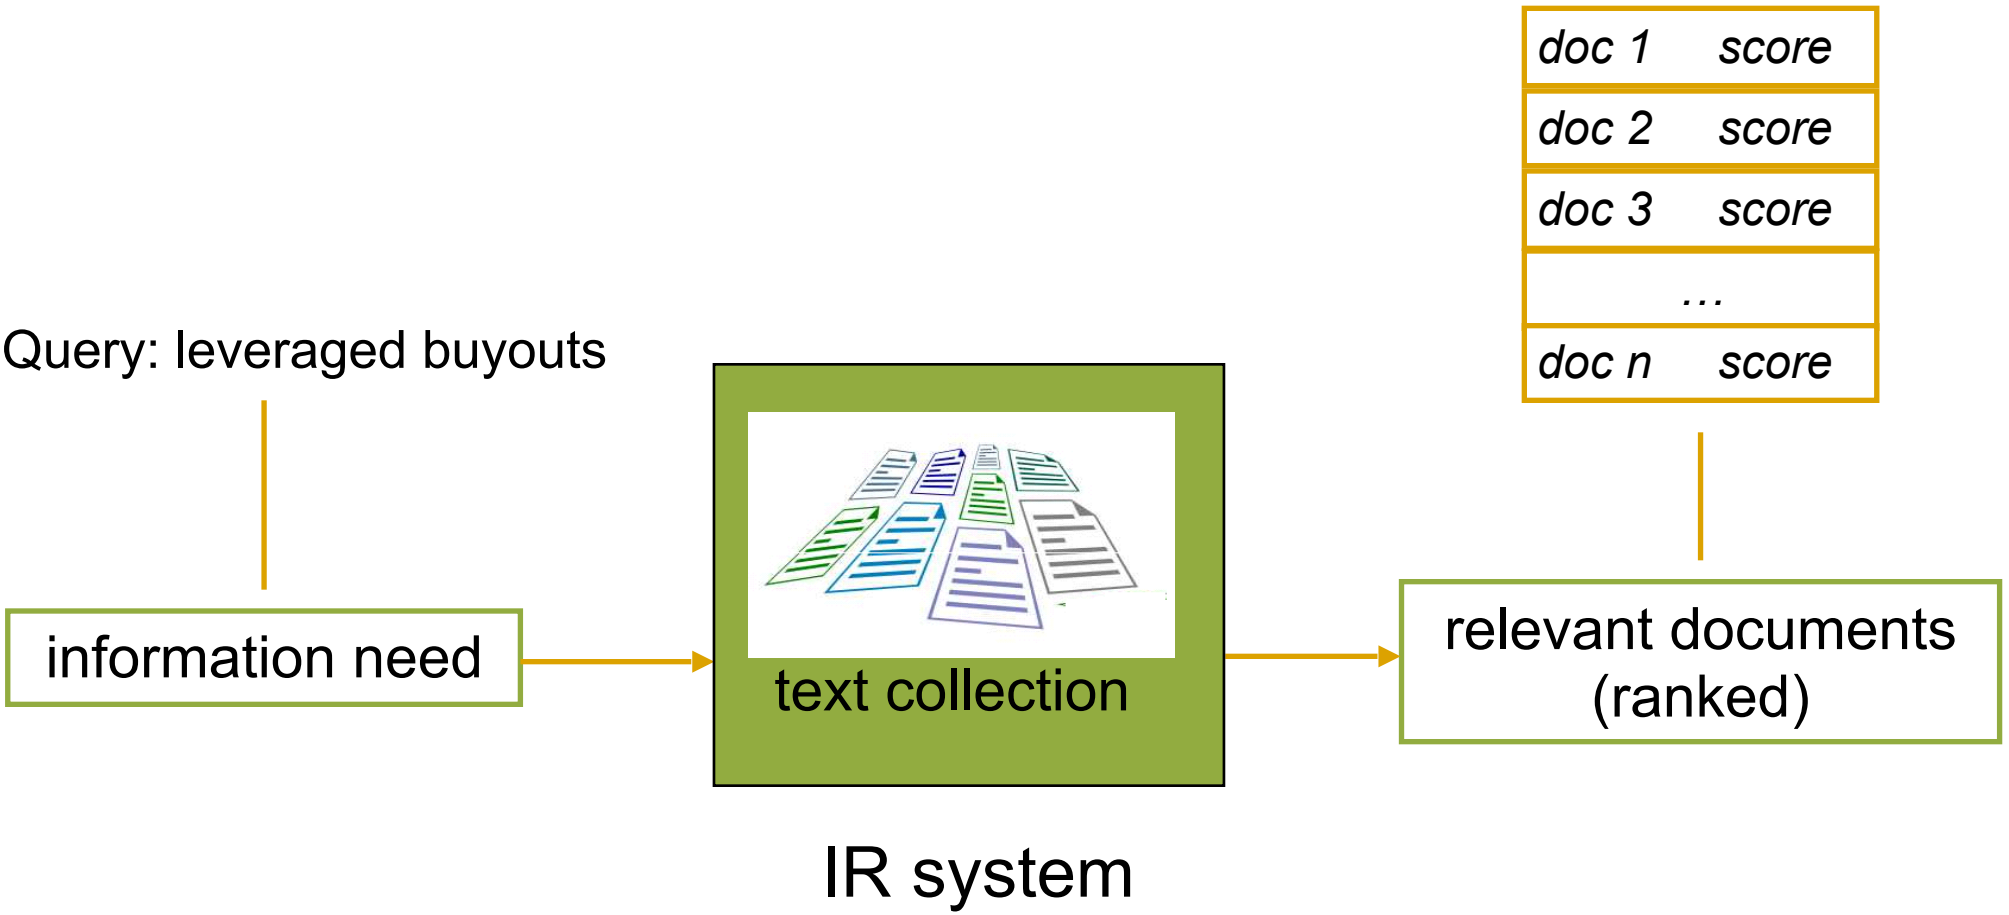

Information retrieval

- Query: *(documents on) leveraged buyouts*

- Query: *(documents on) leveraged buyouts involving more than 100 million dollars that were attempted but failed during 1986 and 1990*

- *I see what I eat = I eat what I see*

  [Mad Hatter, *Alice in Wonderland*]

Question answering (QA)

- Task
  - How many calories are there in a Big Mac?
  - Who is the voice of Miss Piggy?
  - Who was the first American in space?
  - Retrieve not just relevant documents, but return the answer

IBM's Watson

http://www.nytimes.com/video/magazine/124746805578 4/how-does-watson-work.html

Machine translation

- one of the first applications envisioned for NLP techniques
  - *The spirit is willing, but the flesh is weak.*
  - “open”

Dialogue-based systems

- Assistant: Can I help you?
- Customer: I was wondering whether you have any switched brass lampholders.
- Assistant: The brass lampholders are out of stock, but they should be in on Wednesday. The plastic ones are over here...

What kind of NLP system might you be interested in having/building?

?

Topics for Today

- What is NLP?
- **What's involved? Why is it hard?**
- Course structure and requirements

What is the biggest reason that NLP is hard?

(a) "noisy" text

(b) lack of scalability of algorithms

(c) inherent ambiguity of natural language

(d) each natural language is very different from another

(e) [some other reason]

Why is dealing with NL hard?

Ambiguity!!!! …at **all** levels of analysis ☹

- Phonetics and phonology
  - Concerns how words are related to the sounds that realize them. Important for speech-based systems.
    - » "I scream" vs. "ice cream"
    - » "nominal egg"
    - » "It's very hard to wreck a nice beach." (i.e., "It's very hard to recognize speech.")

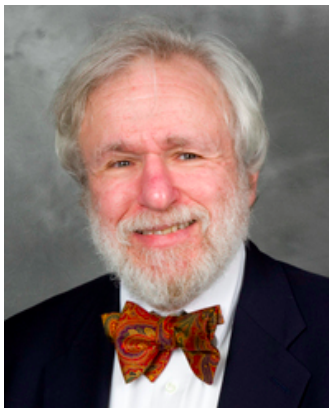

Why is dealing with NL hard?

Ambiguity!!!! …at **all** levels of analysis ☹

- Syntax
  - Concerns sentence structure
  - Different syntactic structure implies different interpretation
    - » Squad helps dog bite victim.
      - [$_{np}$ squad] [$_{vp}$ helps [$_{np}$ dog bite victim]]
      - [$_{np}$ squad] [$_{vp}$ helps [$_{np}$ dog] [$_{inf\text{-}clause}$ bite victim]]
    - » Helicopter powered by human flies.

Why is dealing with NL hard?

Ambiguity!!!! …at **all** levels of analysis ☹

- Semantics
  - Concerns what words mean and how these meanings combine to form sentence meanings.
    - » Red-hot star to wed astronomer.
    - » The once-sagging cloth diaper industry was saved by full dumps.

Why is dealing with NL hard?

Ambiguity!!!! …at **all** levels of analysis ☹

- Discourse
  - Concerns how the immediately preceding sentences affect the interpretation of the next sentence
    - Jack drank the wine on the table. ***It*** was brown and round.
    - Jack saw Sam arrive at the party. Then ***he*** went back to the kitchen to get some chips.
    - Jack saw Sam arrive at the party. ***He*** clearly had drunk too much.

[Adapted from Wilks (1975)]

Why is dealing with NL hard?

Ambiguity!!!! …at **all** levels of analysis ☹

- Pragmatics
  - Concerns how sentences are used in different situations and how use affects the interpretation of the sentence.

    "I just came from Collegetown Bagels."

    - Do you want to go to Collegetown Bagels?
    - Do you want to go to Gimme Coffee?
    - Boy, you look tired.

What topics can we cover?

Language modeling
Phonetic analysis
Morphological analysis
Word-sense disambiguation
Part-of-speech tagging
Parsing
Grammar induction
Semantic analysis
Pronoun resolution
Coreference analysis
NL Generation
Machine translation
Dialogue systems
Information extraction
Information retrieval models
QA systems
Topic models

Reference Material

- Required text book:
  - Jurafsky and Martin, *Speech and Language Processing*, Prentice-Hall, **2nd edition**.

- Other useful references:
  - Manning and Schutze. *Foundations of Statistical NLP*, MIT Press, 1999.
  - Others listed on course web page…

Prereqs, Coursework and Grading

- Prerequisites
  - CS 2110.
- Grading
  - 55%: 3-4 programming projects with short (5-6pg)reports
  - 15%: ~5 critiques of selected research papers
  - 4%: in-class peer-graded midterm
  - 20%: final exam
  - 5%: participation
    You'll be expected to participate in class discussion and class exercises or otherwise demonstrate an interest in the material studied in the course.
  - 1%: course evaluation completion

http://www.cs.cornell.edu/courses/cs4740/